

1. Assume that there are two urns. The first urn contains 4 red balls, 3 blue balls, and 3 white balls. The second urn contains 2 red balls, 4 blue balls, and 4 white balls. You randomly select an urn and take two balls from the urn. The probability that you pick the first urn is 40%. What is the probability that
 - (a) the two balls are red?
 - (b) the second ball is blue?
 - (c) the second ball is blue given that the first ball is red?
2. A fair coin is tossed four times. Let Z be the number of coins that turn up heads.
 - (a) What is the expected value of Z ?
 - (b) What is the variance of Z ?
3. Generally, students who study regularly have a probability of 85% to get an A, 10% to get a B, and 5% to get a C or lower. Unfortunately, some students only study right before a test. These students have a probability of 50% to get an A, 15% to get a B, and 35% to get a C or lower. Thankfully, the students who study regularly make up 90% of the student body.
 - (a) Given that a student received an A on the test, what is the probability that they studied?
 - (b) Given that a student received a B or lower, what is the probability that they didn't study?
4. Let X and Y be random variables. Let $\mathbb{E}[X]$ and $\text{var}[X]$ be the expected value and variance, respectively, of a random variable X .
 - (a) Show that $\mathbb{E}[X + Y] = \mathbb{E}[X] + \mathbb{E}[Y]$.
 - (b) If X and Y are independent, show that $\text{var}[X + Y] = \text{var}[X] + \text{var}[Y]$.
5. Suppose that you are waiting at a bus stop. The waiting time until a bus arrives is T where T is an exponentially distributed random variable with parameter λ i.e. $\Pr(T = t) = \lambda * e^{-\lambda * t} \forall t \geq 0$.
 - (a) Given that you have already waited r seconds, what is the probability that the bus will arrive in d more seconds?
 - (b) What is the average waiting time for the bus i.e. the expected value of T ? Hint: Recall that one way to solve $\int u dv$ is by integration by parts.

6. In this exercise, you will implement the perceptron algorithm. You will be provided with 3 datasets: *data1.csv*, *data2.csv*, and *data3.csv*. Each dataset will have three columns. The first two columns are the attributes of the datapoint and the third column is the label for each datapoint. The attributes have been normalized so that $\|x\| \leq 1$. Each label is either 1 or -1 .

- (a) Plot all datasets. Which datasets are linearly separable?
- (b) Implement the perceptron algorithm as shown in chapter 4 of *A Course in Machine Learning*. To allow for the same results, initialize the hyperplane parameters as 0, iterate through data points in the order provided. Set the maximum iteration number to 1000. For each dataset, provide the hyperplane parameters that are learned by the perceptron algorithm (w and b) and report the total number of updates performed (u). In addition, provide a plot that shows both the hyperplane and the data for each dataset. Based on the total number of updates performed, comment on the convergence of perceptron algorithm for each data set.
- (c) Now, you will compare the rate of convergence for the linearly separable datasets. Recall that the margin $\gamma_{w,b}$ is the distance between the hyperplane defined by $\{w, b\}$ and the nearest point of a set. The margin γ of a set is the largest $\gamma_{w,b}$ for all hyperplanes $\{w, b\}$ that separate the set. As shown in lecture, the number of updates needed to converge is upper bounded by $\frac{1}{\gamma^2}$. Unfortunately, we currently do not have the tools to find γ (will be discussed when the course reaches SVMs). Fortunately, we can use the hyperplane found by the perceptron to get an lower bound on the margin since by definition $\gamma_{w,b} \leq \gamma$ which implies that $\frac{1}{\gamma^2} \leq \frac{1}{\gamma_{w,b}^2}$.
For each linearly separable dataset, calculate the margin $\gamma_{w,b}$ using the learned parameters and compare the theoretical bound to the number of updates that you actually had.